

# Development Individual Project: Presentation

## Developing a Neural Network with the CIFAR-10 image dataset

William Bolton

[w.s.bolton@leeds.ac.uk](mailto:w.s.bolton@leeds.ac.uk)

24<sup>th</sup> May 2024



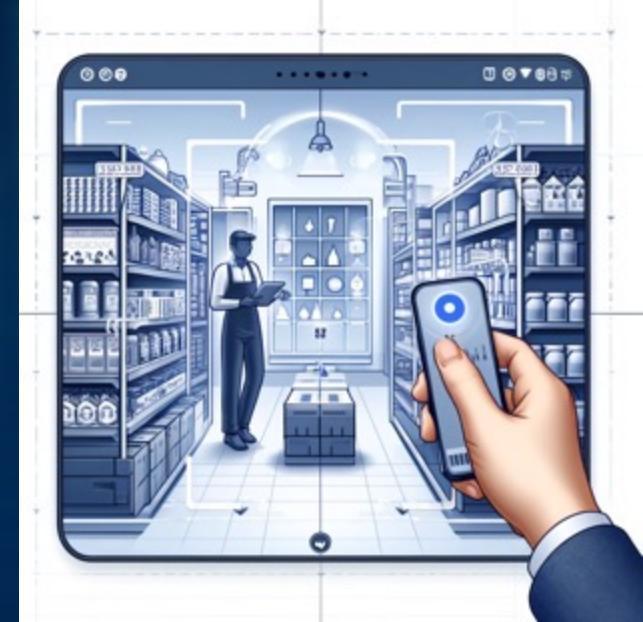
# AI and Object Recognition



Autonomous Vehicles



Diagnostic medical imaging

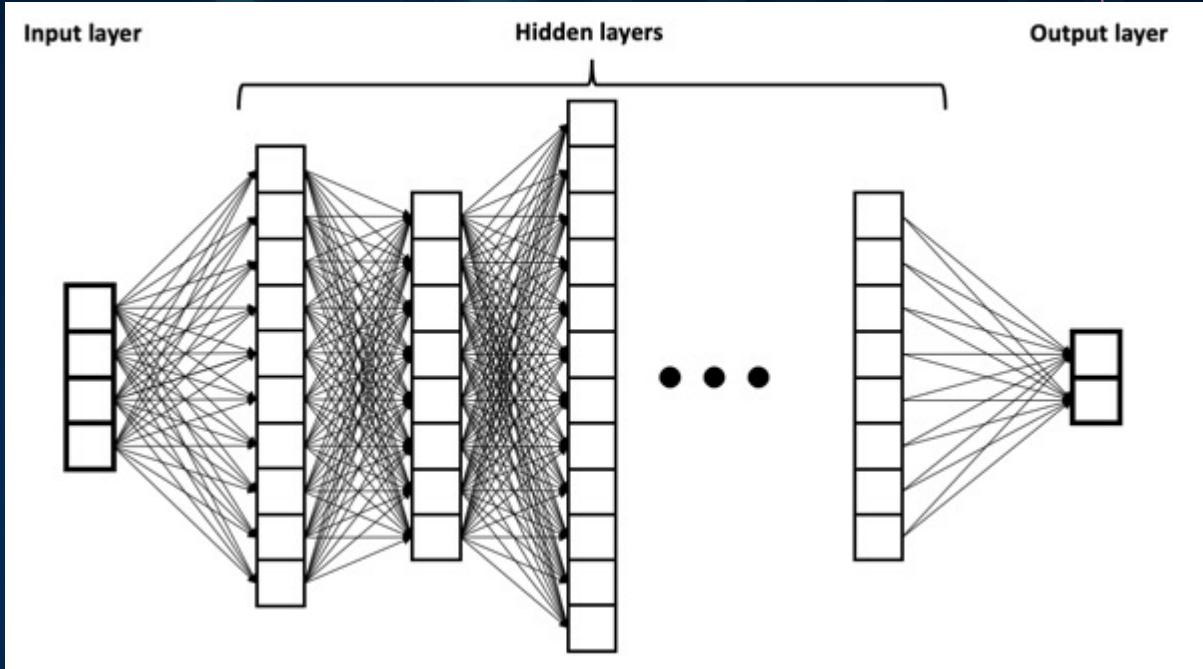
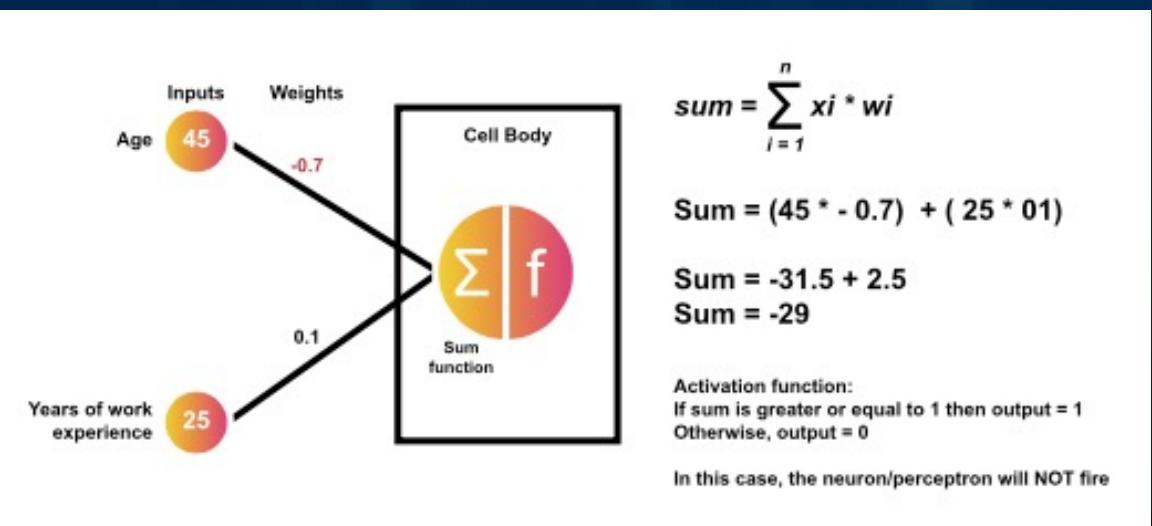


Inventory management

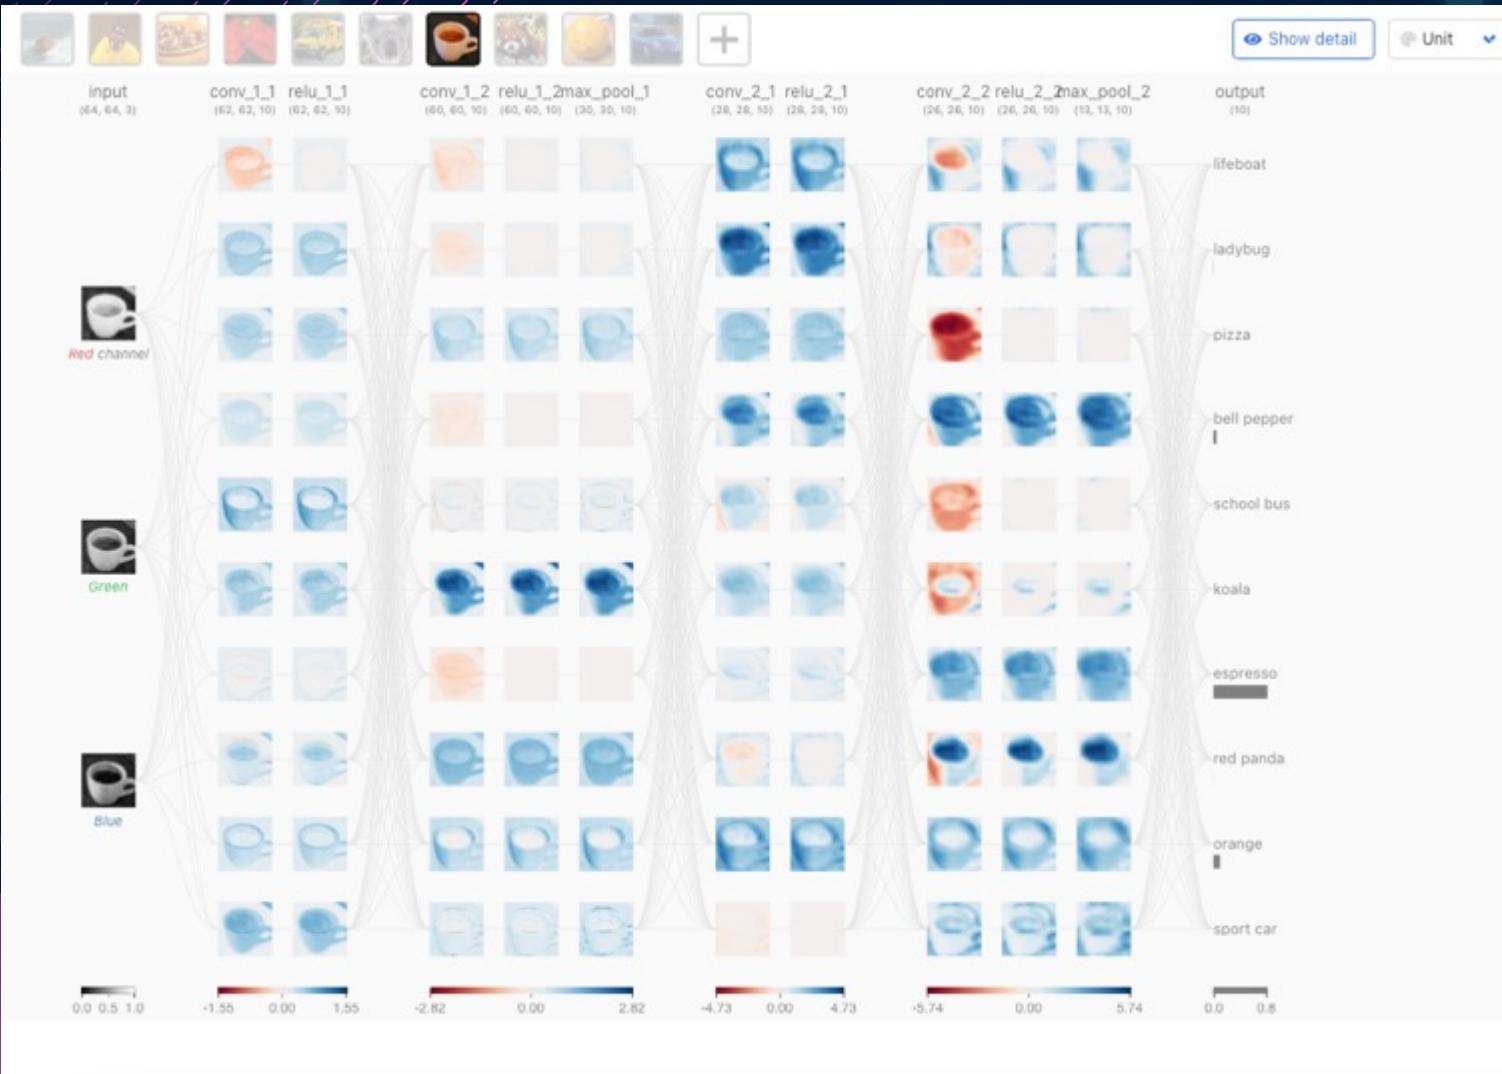
**Object recognition is a critical function of many AI systems that typically required human intelligence.**



# Neural Networks: A backbone of AI



# Architecture of the Artificial Neural Network



- Standard feed-forward convolutional neural network
- Sequential model with Conv2D, MaxPooling2D, Flatten, Dense layers
- Use ReLU activation for hidden layers, softmax for output layer
- Before I could test this design, I prepared and partitioned the data



# Imports, setup, and preparing the dataset

```
1 import os  
2 import pickle  
3 import tensorflow as tf  
4 from tensorflow.keras.datasets import cifar10  
5 from tensorflow.keras.models import Sequential  
6 from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout, MaxPooling2D  
7 from tensorflow.keras.utils import to_categorical  
8 import numpy as np  
9 import matplotlib.pyplot as plt  
10 from sklearn.metrics import precision_score, recall_score, f1_score
```

```
41  
42 # Normalize the data  
43 train_images, test_images = train_images / 255.0, test_images / 255.0  
44
```



# Partitioning the validation set from training data and its importance

```
58 # Ensures representative splitting
59 def representative_split(X, y, test_size=0.2, random_state=42):
60     np.random.seed(random_state)
61     indices = np.arange(X.shape[0])
62     np.random.shuffle(indices)
63     split = int(X.shape[0] * (1 - test_size))
64     train_idx, test_idx = indices[:split], indices[split:]
65     return X[train_idx], X[test_idx], y[train_idx], y[test_idx]
66
67 # Partition the training data to create a validation set using representative split
68 train_images, validation_images, train_labels, validation_labels = representative_split(train_images, train_labels, test_size=0.1)
69
```

```
70 # Data specifics (metadata)
71 print(f"Training samples: {train_images.shape[0]}")
72 print(f"Validation samples: {validation_images.shape[0]}")
73 print(f"Test samples: {test_images.shape[0]}")
```

Training samples: 45000  
Validation samples: 5000  
Test samples: 10000



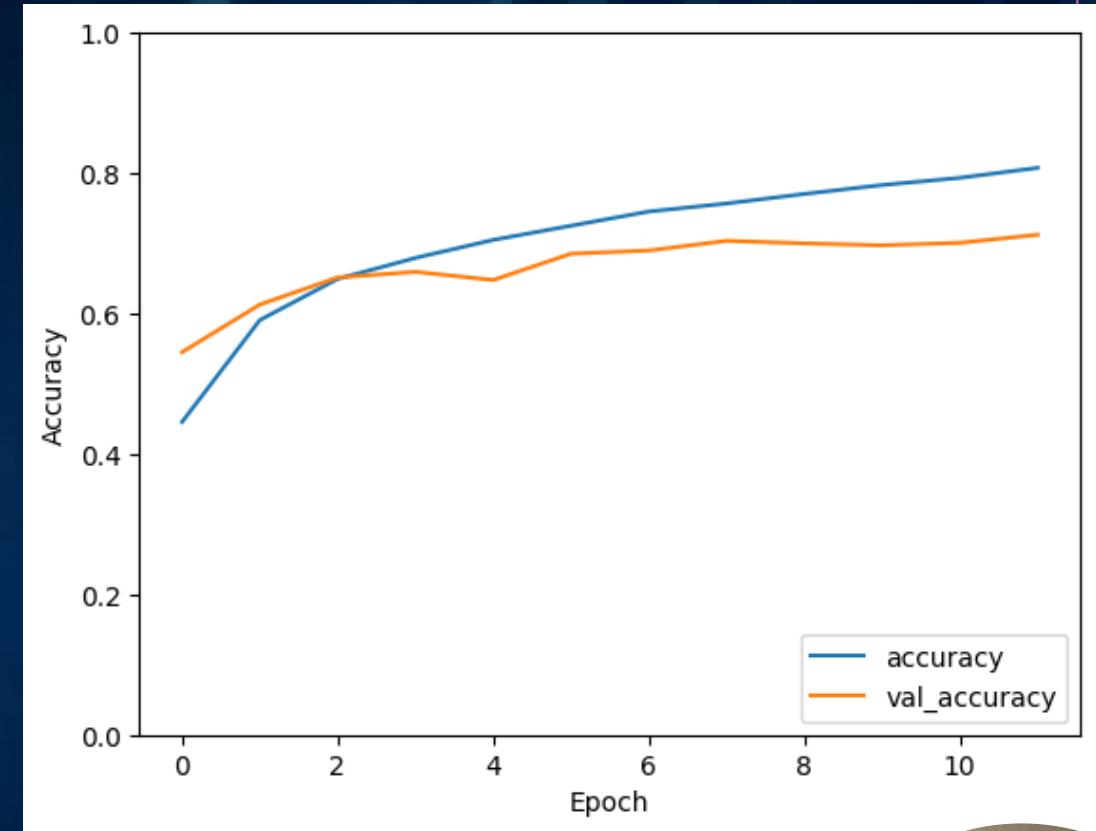
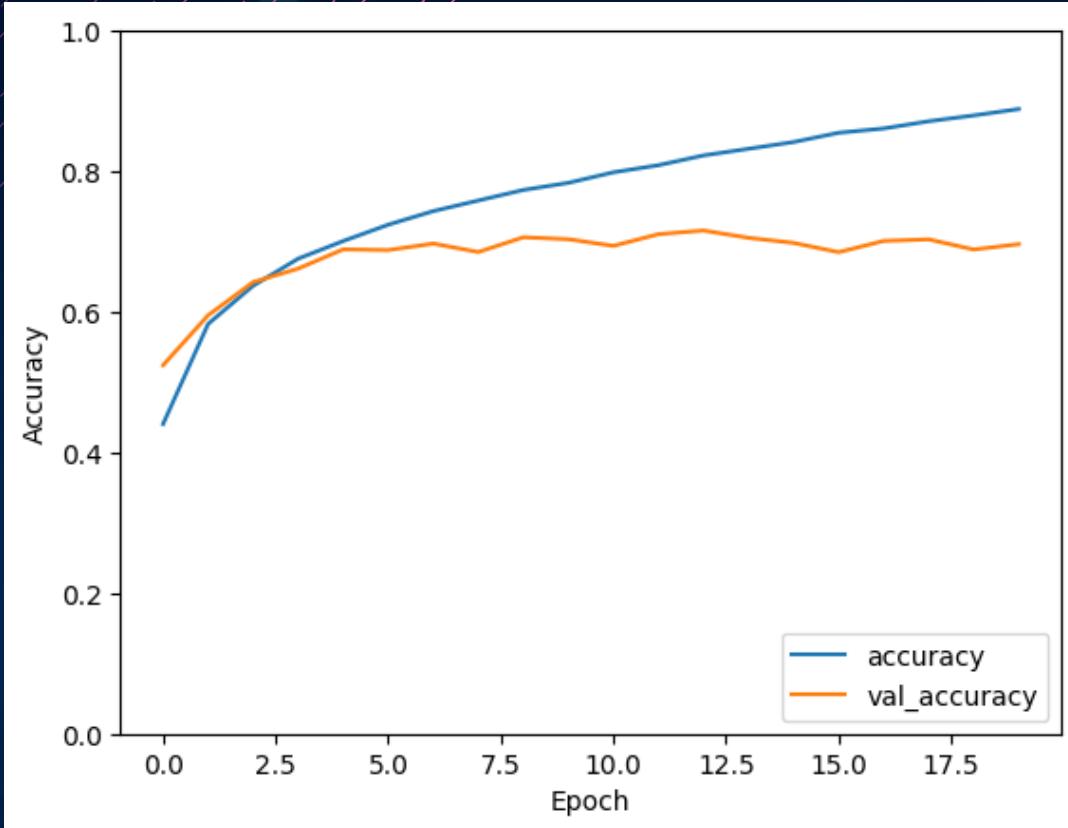
# Activation and Loss Function

$\text{ReLU}(x) = \max(0, x)$

$$L = \frac{1}{N} \sum_{i=1}^N L_i = \frac{1}{N} \sum_{i=1}^N -\log(\hat{y}_{i,y_i})$$



# Epochs in the modelling process



# Neural Network's Design Strategy

```
# Define the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax')
])
```

```
# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=12,
                     validation_data=(validation_images, validation_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(f"Test accuracy: {test_acc}")
```



# Reflecting and conclusions on learnings acquired

```
Epoch 1/12
1407/1407 32s 22ms/step - accuracy: 0.3408 - loss: 1.7885 - val_accuracy: 0.5104 - val_loss: 1.3531
Epoch 2/12
1407/1407 33s 23ms/step - accuracy: 0.5579 - loss: 1.2339 - val_accuracy: 0.6192 - val_loss: 1.0882
Epoch 3/12
1407/1407 34s 24ms/step - accuracy: 0.6261 - loss: 1.0543 - val_accuracy: 0.6406 - val_loss: 1.0188
Epoch 4/12
1407/1407 33s 24ms/step - accuracy: 0.6670 - loss: 0.9489 - val_accuracy: 0.6514 - val_loss: 1.0065
Epoch 5/12
1407/1407 34s 24ms/step - accuracy: 0.6938 - loss: 0.8753 - val_accuracy: 0.6794 - val_loss: 0.9024
Epoch 6/12
1407/1407 34s 24ms/step - accuracy: 0.7166 - loss: 0.8070 - val_accuracy: 0.6838 - val_loss: 0.9094
Epoch 7/12
1407/1407 33s 24ms/step - accuracy: 0.7347 - loss: 0.7534 - val_accuracy: 0.6812 - val_loss: 0.9259
Epoch 8/12
1407/1407 34s 24ms/step - accuracy: 0.7484 - loss: 0.7126 - val_accuracy: 0.7008 - val_loss: 0.8817
Epoch 9/12
1407/1407 33s 23ms/step - accuracy: 0.7608 - loss: 0.6768 - val_accuracy: 0.6938 - val_loss: 0.9330
Epoch 10/12
1407/1407 32s 23ms/step - accuracy: 0.7803 - loss: 0.6331 - val_accuracy: 0.6956 - val_loss: 0.9164
Epoch 11/12
1407/1407 32s 23ms/step - accuracy: 0.7843 - loss: 0.6081 - val_accuracy: 0.7088 - val_loss: 0.8630
Epoch 12/12
1407/1407 32s 22ms/step - accuracy: 0.7958 - loss: 0.5683 - val_accuracy: 0.7074 - val_loss: 0.8699
313/313 - 2s - 8ms/step - accuracy: 0.7061 - loss: 0.8953
Test accuracy: 0.7060999870300293
313/313 3s 8ms/step
Test precision: 0.7108038470261907
Test recall: 0.7061
Test F1 score: 0.7060825735915871
```



# References

CNN Explainer: Learn Convolutional Neural Network Webpage by Wang, J; Turko: R; Shaikh, O; Park, H; Das, N; Hohman, F; Kahng, M; Chau, P. Online at: <https://poloclub.github.io/cnn-explainer/> Accessed 24.05.2024.

Kubat, M., 2017. An introduction to machine learning. Springer.

Learning Multiple Layers of Features from Tiny Images, Alex Krizhevsky, 2009.

File: Example of a deep neural network.png, From Wikimedia Commons, the free media repository. Online at:

[https://commons.wikimedia.org/wiki/File:Example\\_of\\_a\\_deep\\_neural\\_network.png](https://commons.wikimedia.org/wiki/File:Example_of_a_deep_neural_network.png) accessed 24.05.2024.

The Essential Guide to Neural Network Architectures, by Pragati Baheti (Microsoft), V7. Online at: <https://www.v7labs.com/blog/neural-network-architectures-guide> Accessed 24.05.2024.

Techniques for training large neural networks, OpenAI. Online at:  
<https://openai.com/index/techniques-for-training-large-neural-networks/> Accessed 24.05.2024.

Will Cukierski. (2013). CIFAR-10 - Object Recognition in Images. Kaggle.  
<https://kaggle.com/competitions/cifar-10>



# Development Individual Project: Presentation

## Developing a Neural Network with the CIFAR-10 image dataset

William Bolton

[w.s.bolton@leeds.ac.uk](mailto:w.s.bolton@leeds.ac.uk)

24<sup>th</sup> May 2024

